

## Objectives

- Ideas and skills
  - A directory is list of files
  - How to read a directory
  - Types of files and how to determine the type of a file
  - Properties of files and how to determine properties of a file
  - Bit sets and bit masks
  - User and group ID numbers and the passwd database
- System calls and functions
  - opendir, readdir, closedir, seekdir
  - stat, rename
  - chmod, chown, utime
- Commands
  - ls

## Question 1: What does ls do?

ls lists the names of files and reports file attributes

```
[wch@localhost chapter3]$ ls
```

```
chapter3.ppt          slides_chapter3.dvi  slides_chapter3.tex
slides_chapter3.aux    slides_chapter3.log  slides_chapter3.tex~
slides_chapter3.bbl    slides_chapter3.pdf  slides_chapter3.tex.backup
slides_chapter3.blg    slides_chapter3.ps
```

```
[wch@localhost chapter3]$ ls -l
```

```
total 204
```

```
-rw-r--r--  1 wch users 27136 Feb 27 20:13 chapter3.ppt
-rw-r--r--  1 wch users   128 Feb 27 20:36 slides_chapter3.aux
-rw-r--r--  1 wch users     0 Feb 27 20:36 slides_chapter3.bbl
-rw-r--r--  1 wch users   970 Feb 27 20:36 slides_chapter3.blg
-rw-r--r--  1 wch users  3136 Feb 27 20:36 slides_chapter3.dvi
-rw-r--r--  1 wch users 10677 Feb 27 20:36 slides_chapter3.log
-rw-r--r--  1 wch users 14719 Feb 27 20:36 slides_chapter3.pdf
-rw-r--r--  1 wch users 74053 Feb 27 20:36 slides_chapter3.ps
-rw-r--r--  1 wch users  2518 Feb 27 20:36 slides_chapter3.tex
-rw-r--r--  1 wch users  2517 Feb 27 20:36 slides_chapter3.tex~
-rw-r--r--  1 wch users  2181 Feb 27 20:28 slides_chapter3.tex.backup
```

## Question 1: What does ls do?

ls lists other directories, report on other files

-----  
Asking ls about other directories and files  
-----

Example

Action

-----  
ls /tmp                      list name of files in /tmp directory  
ls -l docs                  show attributes of files in docs directory  
ls -l ../Makefile          show attributes of ../Makefile  
ls \*.c                      list names of files matching pattern \*.c  
-----

## Question 2: How does ls work?

A initial guess can be (written in pseudocode):

```
Open directory
while the end of directory is not reached
    display file information
close directory
```

But how to open and read a directory? Any difference from reading a file?

**Facts** A directory is a kind of file that contains a list of names of files and directories. It consists of a sequence of records, each record represents one item, a single file or a single directory.

Every diretory will contain at least the following two items

- . the current directory
- .. the directory one level up

## How to read directory files?

Search in the manual to get some information relevant to directories.

```
$ man -k direct | grep read
```

```
pax          (1)  - read and write file archives and copy director
```

```
readdir      (2)  - read directory entry
```

```
readdir      (3)  - read a directory
```

```
seekdir      (3)  - set the position of the next readdir() call in
```

## man 3 readdir

REaddir(3)

Linux Programmers Manual

REaddir(3)

### NAME

`readdir` - read a directory

### SYNOPSIS

```
#include <sys/types.h>
```

```
#include <dirent.h>
```

```
struct dirent *readdir(DIR *dir);
```

### DESCRIPTION

The `readdir()` function returns a pointer to a `dirent` structure representing the next directory entry in the directory stream pointed to by `dir`. It returns `NULL` on reaching the end-of-file or if an error occurred.

According to POSIX, the `dirent` structure contains a field `char d_name[]` of unspecified size, with at most `NAME_MAX` characters preceding the terminating null character. Use of other fields will harm the portability of your programs. POSIX 1003.1-2001 also documents the field `ino_t d_ino` as an XSI extension.

The data returned by `readdir()` may be overwritten by subsequent calls to `readdir()` for the same directory stream.

### RETURN VALUE

The `readdir()` function returns a pointer to a `dirent` structure, or `NULL` if an error occurs or end-of-file is reached.

### ERRORS

`EBADF` Invalid directory stream descriptor `dir`.

### CONFORMING TO

SVID 3, BSD 4.3, POSIX 1003.1-2001

### SEE ALSO

`read(2)`, `closedir(3)`, `dirfd(3)`, `opendir(3)`, `rewinddir(3)`, `scandir(3)`, `seekdir(3)`, `tellldir(3)`

## Procedure of reading directory

- `opendir (char *)`  
creates a connection, returns a `DIR *`
- `readdir( DIR *)`  
reads next record, returns a pointer to a struct `dirent`
- `closedir (DIR *)`  
closes a connection



## man dirent.h

<dirent.h>(P)

<dirent.h>(P)

### NAME

dirent.h - format of directory entries

### SYNOPSIS

```
#include <dirent.h>
```

### DESCRIPTION

The internal format of directories is unspecified.

The <dirent.h> header shall define the following type:

DIR     A type representing a directory stream.

It shall also define the structure dirent which shall include the following members:

ino_t	d_ino	File serial number.
-------	-------	---------------------

char	d_name[]	Name of entry.
------	----------	----------------

The type ino\_t shall be defined as described in <sys/types.h> .

The character array `d_name` is of unspecified size, but the number of bytes preceding the terminating null byte shall not exceed `{NAME_MAX}`.

The following shall be declared as functions and may also be defined as macros. Function prototypes shall be provided.

```
int          closedir(DIR *);
DIR          *opendir(const char *);
struct dirent *readdir(DIR *);

int          readdir_r(DIR *restrict, struct dirent *restrict,
                    struct dirent **restrict);

void         rewinddir(DIR *);

void         seekdir(DIR *, long);
long         telldir(DIR *);
```

## dirent structure defined in `/usr/include/dirent.h`

```
struct dirent
{
#ifdef __USE_FILE_OFFSET64
    __ino_t d_ino;
    __off_t d_off;
#else
    __ino64_t d_ino;
    __off64_t d_off;
#endif
    unsigned short int d_reclen;
    unsigned char d_type;
    char d_name[256];
};
```

## Logic for listing a directory

```
main() {  
    opendir  
    while (readdir)  
        print d_name  
    closedir  
}
```

## First version of our ls—ls1

```
/** ls1.c
**  purpose  list contents of directory or directories
**  action   if no args, use .  else list files in args
**/
#include      <stdio.h>
#include      <sys/types.h>
#include      <dirent.h>

void do_ls(char []);

main(int ac, char *av[])
{
    if ( ac == 1 )
        do_ls( "." );
    else
        while ( --ac ){
            printf("%s:\n", *++av );
            do_ls( *av );
        }
}

void do_ls( char dirname[] )
/*
```

```
* list files in directory called dirname
*/
{
    DIR          *dir_ptr;          /* the directory */
    struct dirent *direntp;         /* each entry  */

    if ( ( dir_ptr = opendir( dirname ) ) == NULL )
        fprintf(stderr,"ls1: cannot open %s\n", dirname);
    else
    {
        while ( ( direntp = readdir( dir_ptr ) ) != NULL )
            printf("%s\n", direntp->d_name );
        closedir(dir_ptr);
    }
}
```

## Compile and run our first version of ls–ls1

```
$ cc -o ls1 ls1.c
```

```
$ ls1
```

```
Makefile
```

```
ls1.c
```

```
.
```

```
ls2.c
```

```
fileinfo.c
```

```
filesize.c
```

```
a.out
```

```
..
```

```
$ ls
```

```
a.out  fileinfo.c  filesize.c  ls1.c  ls2.c  Makefile
```

## Comparison between our ls1 and ls

- Not sorted  
we can read all the file names into an array and then use qsort to sort the array
- No column  
read the names into an array and print them out in rows and columns.
- Lists . files  
To fix it, just suppress . and .. names
- No -l info  
To be discussed.



## Information shown by `ls -l`

Each line consists of the following seven fields:

- **mode:** The first character in each line represents the types of the file. The remaining nine characters represent the access permissions (read write execute for user, group, everyone else).
- **owner** Each file belongs to a user. The username of the owner is shown.
- **group:** Each file belongs to a group of users too.
- **size:** The number of bytes in the file
- **last-modified** This field consists of three strings representing the last-modified time. For newer files, the output includes the month, day, and time For older files, `ls` shows the month, day, and year.
- **name:** The name of the file

## How does ls -l work

```
$ man -k file | grep -i status
```

```
fileno [ferror]      (3)- check and reset stream status
```

```
fstat [stat]         (2)- get file status
```

```
ifcfg-ppp0 [pppoe]   (5)- Configuration file used by adsl-start(8),  
                        adsl-stop(8), adsl-status(8) and  
                        adsl-connect(8)
```

```
lam_rfstate          (2)- Report status of remote LAM file  
                        descriptors
```

```
lstat [stat]         (2)- get file status
```

```
stat                 (1)- display file or filesystem status
```

```
stat                 (2)- get file status
```

## How does ls -l work

### NAME

`stat, fstat, lstat` - get file status

### SYNOPSIS

```
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>
```

```
int stat(const char *file_name, struct stat *buf);
int fstat(int filedes, struct stat *buf);
int lstat(const char *file_name, struct stat *buf);
```

### DESCRIPTION

These functions return information about the specified file. You do not need any access rights to the file to get this information but you need search rights to all directories named in the path leading to the file.

`stat` stats the file pointed to by `file_name` and fills in `buf`.

`lstat` is identical to `stat`, except in the case of a symbolic link, where the link itself is stat-ed, not the file that it refers to.

`fstat` is identical to `stat`, only the open file pointed to by `filedes` (as returned by `open(2)`) is stat-ed in place of `file_name`.

They all return a `stat` structure, which contains the following fields:

```
struct stat {
    dev_t      st_dev;      /* device */
    ino_t      st_ino;      /* inode */
    mode_t     st_mode;     /* protection */
    nlink_t    st_nlink;    /* number of hard links */
    uid_t      st_uid;      /* user ID of owner */
    gid_t      st_gid;      /* group ID of owner */
    dev_t      st_rdev;     /* device type (if inode device) */
    off_t      st_size;     /* total size, in bytes */
    blksize_t  st_blksize;  /* blocksize for filesystem I/O */
    blkcnt_t   st_blocks;   /* number of blocks allocated */
    time_t     st_atime;    /* time of last access */
    time_t     st_mtime;    /* time of last modification */
    time_t     st_ctime;    /* time of last status change */
};
```

The value `st_size` gives the size of the file (if it is a regular file or a symlink) in bytes. The size of a symlink is the length of the pathname it contains, without trailing NUL.

The value `st_blocks` gives the size of the file in 512-byte blocks. (This may be smaller than `st_size/512` e.g. when the file has holes.) The value `st_blksize` gives the "preferred" blocksize for efficient file system I/O. (Writing to a file in smaller chunks may cause an inefficient read-modify-rewrite.)

Not all of the Linux filesystems implement all of the time fields. Some file system types allow mounting in such a way that file accesses do not cause an update of the `st_atime` field. (See `noatime` in `mount(8)`.)

The field `st_atime` is changed by file accesses, e.g. by `execve(2)`, `mknod(2)`, `pipe(2)`, `utime(2)` and `read(2)` (of more than zero bytes). Other routines, like `mmap(2)`, may or may not update `st_atime`.

The field `st_mtime` is changed by file modifications, e.g. by `mknod(2)`, `truncate(2)`, `utime(2)` and `write(2)` (of more than zero bytes). Moreover, `st_mtime` of a directory is changed by the creation or deletion of files in that directory. The `st_mtime` field is not changed for changes in owner, group, hard link count, or mode.

The field `st_ctime` is changed by writing or by setting inode information (i.e., owner, group, link count, mode, etc.).

...

## How does ls -l work

```
/* filesize.c - prints size of passwd file */

#include <stdio.h>
#include <sys/stat.h>

int main()
{
    struct stat infobuf;                /* place to store info */

    if ( stat( "/etc/passwd", &infobuf) == -1 ) /* get info */
        perror("/etc/passwd");
    else
        printf(" The size of /etc/passwd is %d\n", infobuf.st_size );
}
```

## How does ls -l work

```
/* statinfo.c - demonstrates using stat() to obtain
 *             file information.
 *             - some members are just numbers...
 */
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>

void show_stat_info(char *, struct stat *);

int main(int ac, char *av[])
{
    struct stat info;          /* buffer for file info */

    if (ac>1)
        if( stat(av[1], &info) != -1 ){
            show_stat_info( av[1], &info );
            return 0;
        }
        else
            perror(av[1]); /* report stat() errors */
    return 1;
}
```

```
void show_stat_info(char *fname, struct stat *buf)
/*
 * displays some info from stat in a name=value format
 */
{
    printf("    mode: %o\n", buf->st_mode);           /* type + mode */
    printf("  links: %d\n", buf->st_nlink);          /* # links      */
    printf("   user: %d\n", buf->st_uid);            /* user id      */
    printf("  group: %d\n", buf->st_gid);            /* group id     */
    printf("   size: %d\n", buf->st_size);           /* file size    */
    printf("modtime: %d\n", buf->st_mtime);          /* modified     */
    printf("   name: %s\n", fname );                /* filename     */
}
```



## How does `ls -l` work

```
$ ./fileinfo fileinfo.c
```

```
mode: 100755
```

```
links: 1
```

```
user: 500
```

```
group: 500
```

```
size: 1152
```

```
modtime: 1124320515
```

```
name: fileinfo.c
```

```
$ ls -l fileinfo.c
```

```
-rwxr-xr-x  1 wch wch 1152 Aug 17  2005 fileinfo.c
```

There are several differences between the output of our `fileinfo` and `ls -l`: mode, user, group, modtime. For modtime, we can use `ctime` to convert it to string.

## Converting file mode to a string




Figure 1: File type and access coding

- The first four bits represent the type of the file. Four bits can contain 16 possible patterns of 1s and 0s. Each pattern can represent one file type.
- The next three bits are for special attributes of a file. Each bit corresponds to a special attribute.

## Converting file mode to a string

We use mask to zero out all the bits except the bits we are interested in.

For example  $011001010101 \ \&\& \ 111000000000 = 011000000000$

Here 111000000000 is the mask.

## Using masking to decode file types

The definition of bits pattern for file types are in  
`/usr/include/bits/stat.h`

```
/* Encoding of the file mode.  */

#define __S_IFMT 0170000 /* These bits determine file type.  */

/* File types.  */
#define __S_IFDIR 0040000 /* Directory.  */
#define __S_IFCHR 0020000 /* Character device.  */
#define __S_IFBLK 0060000 /* Block device.  */
#define __S_IFREG 0100000 /* Regular file.  */
#define __S_IFIFO 0010000 /* FIFO.  */
#define __S_IFLNK 0120000 /* Symbolic link.  */
#define __S_IFSOCK 0140000 /* Socket.  */
```

## Using masking to decode file types

The macros to check the file types is in `/usr/include/sys/stat.h`

```
/* Test macros for file types. */
```

```
#define __S_ISTYPE(mode, mask) (((mode) & __S_IFMT) == (mask))
```

```
#define S_ISDIR(mode)    __S_ISTYPE((mode), __S_IFDIR)
```

```
#define S_ISCHR(mode)    __S_ISTYPE((mode), __S_IFCHR)
```

```
#define S_ISBLK(mode)    __S_ISTYPE((mode), __S_IFBLK)
```

```
#define S_ISREG(mode)    __S_ISTYPE((mode), __S_IFREG)
```

```
#ifdef __S_IFIFO
```

```
# define S_ISFIFO(mode)  __S_ISTYPE((mode), __S_IFIFO)
```

```
#endif
```

```
#ifdef __S_IFLNK
```

```
# define S_ISLNK(mode)   __S_ISTYPE((mode), __S_IFLNK)
```

```
#endif
```

## Using masking to decode file types

The bit pattern for user access permission is defined in `/usr/include/bits/stat.h`

```
/* Protection bits.  */
#define __S_ISUID 04000 /* Set user ID on execution.  */
#define __S_ISGID 02000 /* Set group ID on execution.  */
#define __S_ISVTX 01000 /* Save swapped text after use (sticky).  */
#define __S_IREAD 0400 /* Read by owner.  */
#define __S_IWRITE 0200 /* Write by owner.  */
#define __S_IEXEC 0100 /* Execute by owner.  */
```

The bit pattern for group and other permission is defined in `/usr/include/sys/stat.h`, which includes `/usr/include/bits/stat.h`

```
#define S_IRGRP (S_IRUSR >> 3) /* Read by group.  */
#define S_IWGRP (S_IWUSR >> 3) /* Write by group.  */
#define S_IXGRP (S_IXUSR >> 3) /* Execute by group.  */
/* Read, write, and execute by group.  */
#define S_IRWXG (S_IRWXU >> 3)
#define S_IROTH (S_IRGRP >> 3) /* Read by others.  */
#define S_IWOTH (S_IWGRP >> 3) /* Write by others.  */
#define S_IXOTH (S_IXGRP >> 3) /* Execute by others.  */
/* Read, write, and execute by others.  */
#define S_IRWXO (S_IRWXG >> 3)
```

## Using masking to decode file types

```
void mode_to_letters( int mode, char str[] )
{
    strcpy( str, "-----" );          /* default=no perms */

    if ( S_ISDIR(mode) ) str[0] = 'd';  /* directory?      */
    if ( S_ISCHR(mode) ) str[0] = 'c';  /* char devices    */
    if ( S_ISBLK(mode) ) str[0] = 'b';  /* block device    */

    if ( mode & S_IRUSR ) str[1] = 'r';  /* 3 bits for user */
    if ( mode & S_IWUSR ) str[2] = 'w';
    if ( mode & S_IXUSR ) str[3] = 'x';

    if ( mode & S_IRGRP ) str[4] = 'r';  /* 3 bits for group */
    if ( mode & S_IWGRP ) str[5] = 'w';
    if ( mode & S_IXGRP ) str[6] = 'x';

    if ( mode & S_IROTH ) str[7] = 'r';  /* 3 bits for other */
    if ( mode & S_IWOTH ) str[8] = 'w';
    if ( mode & S_IXOTH ) str[9] = 'x';
}
```

## Converting User ID to Strings

```
root:x:0:0:root:/root:/bin/bash
bin:x:1:1:bin:/bin:/sbin/nologin
daemon:x:2:2:daemon:/sbin:/sbin/nologin
adm:x:3:4:adm:/var/adm:/sbin/nologin
lp:x:4:7:lp:/var/spool/lpd:/sbin/nologin
sync:x:5:0:sync:/sbin:/bin/sync
shutdown:x:6:0:shutdown:/sbin:/sbin/shutdown
halt:x:7:0:halt:/sbin:/sbin/halt
mail:x:8:12:mail:/var/spool/mail:/sbin/nologin
news:x:9:13:news:/etc/news:
mysql:x:27:27:MySQL Server:/var/lib/mysql:/bin/bash
wch:x:500:100:Changhua Wu:/home/wch:/bin/bash
tomcat:x:91:91:Tomcat:/usr/share/tomcat5:/bin/sh
```

The first field is the user name, the second field is the encrypted password, the third field is the user id, the fourth field is the group id. The other fields show the real name of the user, the home directory, and the default shell.



## Converting User ID to Strings

However, on some Unix systems, `/etc/passwd` does not contain all the users. To ensure portability, we use `getpwuid`.

```
struct passwd *getpwuid(uid_t uid);
```

The `getpwuid()` function returns a pointer to a structure containing the broken out fields of a line from `/etc/passwd` for the entry that matches the user `uid`. The `passwd` structure is defined in `pwd.h` as follows:

```
struct passwd {
    char    *pw_name;        /* user name */
    char    *pw_passwd;      /* user password */
    uid_t   pw_uid;          /* user id */
    gid_t    pw_gid;          /* group id */
    char    *pw_gecos;        /* real name */
    char    *pw_dir;          /* home directory */
    char    *pw_shell;        /* shell program */
};
```

So we can use this function to get user name from user id

```
char * uid_to_name ( uid_t uid) {
    return getpwuid (uid)->pw_name;
}
```

## Converting Group ID to Strings

/etc/group is the list of groups. Like /etc/passwd, it is a plain text file that looks like the following:

```
root:x:0:root
apache:x:48:
iiimd:x:101:
mysql:x:27:
screen:x:84:
users:x:100:wch
wch:x:500:
tomcat:x:91:
```

The name of the group is the first field, the group password is the second field, the group ID is the third field. The fourth field is a comma-separated list of usernames.

## Converting Group ID to Strings

```
struct group *getgrgid(gid_t gid);
```

The `getgrgid()` function returns a pointer to a structure containing the group information from `/etc/group` for the entry that matches the group id.

```
char *gid_to_name ( gid_t uid){  
    return getgrgid(gid)->gr_name;  
}
```

## Our second version of ls—ls2

```
/* ls2.c
 * purpose  list contents of directory or directories
 * action   if no args, use .  else list files in args
 * note     uses stat and pwd.h and grp.h
 * BUG: try ls2 /tmp
 */
#include <stdio.h>
#include <sys/types.h>
#include <dirent.h>
#include <sys/stat.h>

void do_ls(char[]);
void dostat(char *);
void show_file_info( char *, struct stat *);
void mode_to_letters( int , char [] );
char *uid_to_name( uid_t );
char *gid_to_name( gid_t );

main(int ac, char *av[])
{
    if ( ac == 1 )
        do_ls( "." );
    else
```

```
        while ( --ac ){
            printf("%s:\n", *++av );
            do_ls( *av );
        }
}

void do_ls( char dirname[] )
/*
 *   list files in directory called dirname
 */
{
    DIR *dir_ptr;                /* the directory */
    struct dirent *direntp;      /* each entry */

    if ( ( dir_ptr = opendir( dirname ) ) == NULL )
        fprintf(stderr,"ls1: cannot open %s\n", dirname);
    else
    {
        while ( ( direntp = readdir( dir_ptr ) ) != NULL )
            dostat( direntp->d_name );
        closedir(dir_ptr);
    }
}
```

```
void dostat( char *filename )
{
    struct stat info;

    if ( stat(filename, &info) == -1 )           /* cannot stat */
        perror( filename );                     /* say why */
    else                                           /* else show info */
        show_file_info( filename, &info );
}

void show_file_info( char *filename, struct stat *info_p )
/*
 * display the info about 'filename'.  The info is stored in struct at *info_p
 */
{
    char *uid_to_name(), *ctime(), *gid_to_name(), *filemode();
    void mode_to_letters();
    char    modestr[11];

    mode_to_letters( info_p->st_mode, modestr );

    printf( "%s"      , modestr );
    printf( "%4d "    , (int) info_p->st_nlink);
    printf( "%-8s "   , uid_to_name(info_p->st_uid) );
}
```

```
    printf( "%-8s " , gid_to_name(info_p->st_gid) );
    printf( "%8ld " , (long)info_p->st_size);
    printf( "%.12s ", 4+ctime(&info_p->st_mtime));
    printf( "%s\n" , filename );

}
/*
 * utility functions
 */

/*
 * This function takes a mode value and a char array
 * and puts into the char array the file type and the
 * nine letters that correspond to the bits in mode.
 * NOTE: It does not code setuid, setgid, and sticky
 * codes
 */
void mode_to_letters( int mode, char str[] )
{
    strcpy( str, "-----" );          /* default=no perms */

    if ( S_ISDIR(mode) ) str[0] = 'd';  /* directory?      */
    if ( S_ISCHR(mode) ) str[0] = 'c';  /* char devices     */
    if ( S_ISBLK(mode) ) str[0] = 'b';  /* block device     */
}
```

```
    if ( mode & S_IRUSR ) str[1] = 'r';    /* 3 bits for user */
    if ( mode & S_IWUSR ) str[2] = 'w';
    if ( mode & S_IXUSR ) str[3] = 'x';

    if ( mode & S_IRGRP ) str[4] = 'r';    /* 3 bits for group */
    if ( mode & S_IWGRP ) str[5] = 'w';
    if ( mode & S_IXGRP ) str[6] = 'x';

    if ( mode & S_IROTH ) str[7] = 'r';    /* 3 bits for other */
    if ( mode & S_IWOTH ) str[8] = 'w';
    if ( mode & S_IXOTH ) str[9] = 'x';
}

#include <pwd.h>

char *uid_to_name( uid_t uid )
/*
 *   returns pointer to username associated with uid, uses getpw()
 */
{
    struct passwd *getpwuid(), *pw_ptr;
    static char numstr[10];

    if ( ( pw_ptr = getpwuid( uid ) ) == NULL ){
```



```
        sprintf(numstr,"%d", uid);
        return numstr;
    }
    else
        return pw_ptr->pw_name ;
}

#include    <grp.h>

char *gid_to_name( gid_t gid )
/*
 *    returns pointer to group number gid. used getgrgid(3)
 */
{
    struct group *getgrgid(), *grp_ptr;
    static  char numstr[10];

    if ( ( grp_ptr = getgrgid(gid) ) == NULL ){
        sprintf(numstr,"%d", gid);
        return numstr;
    }
    else
        return grp_ptr->gr_name;
}
```

## Three special bits

This three bits is defined in `/usr/include/bits/stat.h`

- Set-User-ID bit: set the effective user ID as the uses ID of the owner  
`#define __S_ISUID 04000 /* Set user ID on execution.*/`
- Set-Group-ID: set the effective group id as the group id of the file  
`#define __S_ISGID 02000 /* Set group ID on execution.*/`
- The Sticky Bit: It has two different uses, one for files and one for directories.  
`#define __S_ISVTX 01000 /* Save swapped text after use (sticky).*/`

A quite detailed description of the sticky bit for files and directories.

<http://www.faqs.org/faqs/hp/hpux-faq/section-70.html>

## Setting and modifying the properties of a file

- Type of a file

The type of file is determined at the time of creation. It is not possible to change the type of a file.

- Permission bits and special bits

- Establishing the mode of a file

`fd = creat("newfile", 0744);` This will create newfile and requests an initial set of permission bits `rwxr-r-`

- changing the Mode of a file

`chmod("/tmp/myfile", 04764);`

or

`chmod("/tmp/myfile", S_ISUID | S_IRWXU | S_IRGRP | S_IWGRP | S_IROTH);`

## man 2 chmod

### NAME

chmod, fchmod - change permissions of a file

### SYNOPSIS

```
#include <sys/types.h>
#include <sys/stat.h>
```

```
int chmod(const char *path, mode_t mode);
int fchmod(int fildes, mode_t mode);
```

### DESCRIPTION

The mode of the file given by path or referenced by fildes is changed.

Modes are specified by oring the following:

S_ISUID	04000	set user ID on execution
S_ISGID	02000	set group ID on execution
S_ISVTX	01000	sticky bit
S_IRUSR	(S_IREAD) 00400	read by owner
S_IWUSR	(S_IWRITE) 00200	write by owner
S_IXUSR	(S_IEXEC) 00100	execute/search by owner
S_IRGRP	00040	read by group
S_IWGRP	00020	write by group

S_IXGRP	00010	execute/search by group
S_IROTH	00004	read by others
S_IWOTH	00002	write by others
S_IXOTH	00001	execute/search by others

The effective UID of the process must be zero or must match the owner of the file.

## Changing the owner and group a file

```
chown ("file1", 200, 40 );
```

```
$man 2 chown
```

### NAME

chown, fchown, lchown - change ownership of a file

### SYNOPSIS

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
int chown(const char *path, uid_t owner, gid_t group);
```

```
int fchown(int fd, uid_t owner, gid_t group);
```

```
int lchown(const char *path, uid_t owner, gid_t group);
```

### DESCRIPTION

The owner of the file specified by path or by fd is changed. Only the super-user may change the owner of a file. The owner of a file may change the group of the file to any group of which that owner is a member. The super-user may change the group arbitrarily.

If the owner or group is specified as -1, then that ID is not changed.

When the owner or group of an executable file are changed by a non-super-user, the S\_ISUID and S\_ISGID mode bits are cleared. POSIX does not specify whether this also should happen when root does the chown; the Linux behaviour depends on the kernel version. In case of a non-group-executable file (with clear S\_IXGRP bit) the S\_ISGID bit indicates mandatory locking, and is not cleared by a chown.

### RETURN VALUE

On success, zero is returned. On error, -1 is returned, and `errno` is set appropriately.

## Changing the time of a file—utime

The utime system call sets the modification and access times for a file.

UTIME(2)

Linux Programmers Manual

UTIME(2)

### NAME

utime, utimes - change access and/or modification times of an inode

### SYNOPSIS

```
#include <sys/types.h>
```

```
#include <utime.h>
```

```
int utime(const char *filename, const struct utimbuf *buf);
```

```
#include <sys/time.h>
```

```
int utimes(char *filename, struct timeval *tvp);
```

### DESCRIPTION

utime changes the access and modification times of the inode specified by filename to the actime and modtime fields of buf respectively. If buf is NULL, then the access and modification times of the file are set to the current time. The utimbuf structure is:



```
struct utimbuf {
    time_t actime; /* access time */
    time_t modtime; /* modification time */
};
```

In the Linux DLL 4.4.1 libraries, `utimes` is just a wrapper for `utime`: `tvp[0].tv_sec` is `actime`, and `tvp[1].tv_sec` is `modtime`. The `timeval` structure is:

```
struct timeval {
    long    tv_sec;          /* seconds */
    long    tv_usec;        /* microseconds */
};
```

### RETURN VALUE

On success, zero is returned. On error, -1 is returned, and `errno` is set appropriately.

## Changing the name of a file—rename

RENAME(2)

Linux Programmers Manual

RENAME(2)

### NAME

rename - change the name or location of a file

### SYNOPSIS

```
#include <stdio.h>
```

```
int rename(const char *oldpath, const char *newpath);
```

### DESCRIPTION

rename renames a file, moving it between directories if required.

Any other hard links to the file (as created using `link(2)`) are unaffected.

If `newpath` already exists it will be atomically replaced (subject to a few conditions - see ERRORS below), so that there is no point at which another process attempting to access `newpath` will find it missing.

If `newpath` exists but the operation fails for some reason rename guarantees to leave an instance of `newpath` in place.

However, when overwriting there will probably be a window in which both oldpath and newpath refer to the file being renamed.

If oldpath refers to a symbolic link the link is renamed; if newpath refers to a symbolic link the link will be overwritten.

### RETURN VALUE

On success, zero is returned. On error, -1 is returned, and errno is set appropriately.