

Objectives

- Ideas and Skills
 - User's view of the Unix file system tree
 - Internal structure of Unix file system: inodes and data blocks
 - How directories are connected
 - Hard links, symbolic links: ideas and system calls
 - How pwd works
 - Mounting file systems
- System calls and functions
 - mkdir, rmdir, chdir
 - link, unlink, rename, symlink
- pwd

Directories and files




Figure 1: A tree of directories

Unix Commands for directories and files

```
$mkdir demodir
```

```
$cd demodir
```

```
$mkdir demodir/a
```

```
$mkdir demodir/c
```

```
$cd c
```

```
mkdir d1 d2
```

```
$cd demodir
```

```
$cp /etc/group x
```

```
$cat x
```

```
root:x:0:root
```

```
bin:x:1:root,bin,daemon
```

```
disk:x:6:roo
```

```
...
```

```
$cp x copy.of.x
```

```
$rm copy.of.x
```

Internal structure of the Unix file system

- Abstraction Zero: From platters to partitions, usually partitions stop at the boundary of cylinders/tracks.
- Abstraction One: From partition to array of blocks.
Index the sector for each partition in some order by a numbering system.




Figure 2: A tree of directories

- Abstraction Two: From an array of blocks to three regions:




Figure 3: A tree of directories

- the **superblock**: the first block in the file system. It contains information about the organization of the file system itself.
- the **inode table**: inodes store the information of files such as size, user ID of the owner, time of last modification.
- the **data area**: the content of a file is in this area. Some files may take more than one block.

Internal structure of the Unix file system

There are four steps in creating a file:

- Store properties: The kernel finds an available inode and write some file properties to the inode.
- Store data: The kernel finds a number of unoccupied/free blocks to save the content of the file.
- Record allocation: write the sequence of blocks taken by the file into the disk allocation section of the inode.
- Add filename into directory: The kernel adds an entry (inode number, filename) to some directory.




Figure 4: Internal structure of a file

Internal structure of the Unix file system

A directory is a special kind of files that contains a list of names of files and their inode numbers.

```
[wch@localhost temp]$ ls -la
1210078 .      1210797 glut-3.6          1211631 svl-1.5
    32705 ..   3679212 glut-3.6.tar.gz  3679211 svl-1.5.tar.gz
```

Multiple files may point to the same inode.

How cat works

\$cat userlist

There are several steps in executing the above command


- Search in the directory for the filename and get the inode number
- Locate and read inode in the inode table of the file system.
- Read the data blocks

```
$ cat userlist
```

123	.
833	..
4004	hello.c
47	userlist

From filename to file contents:

Search the directory for a filename,
find the associated inum,
use the inum to locate the inode,
the inode contains a list of data blocks.

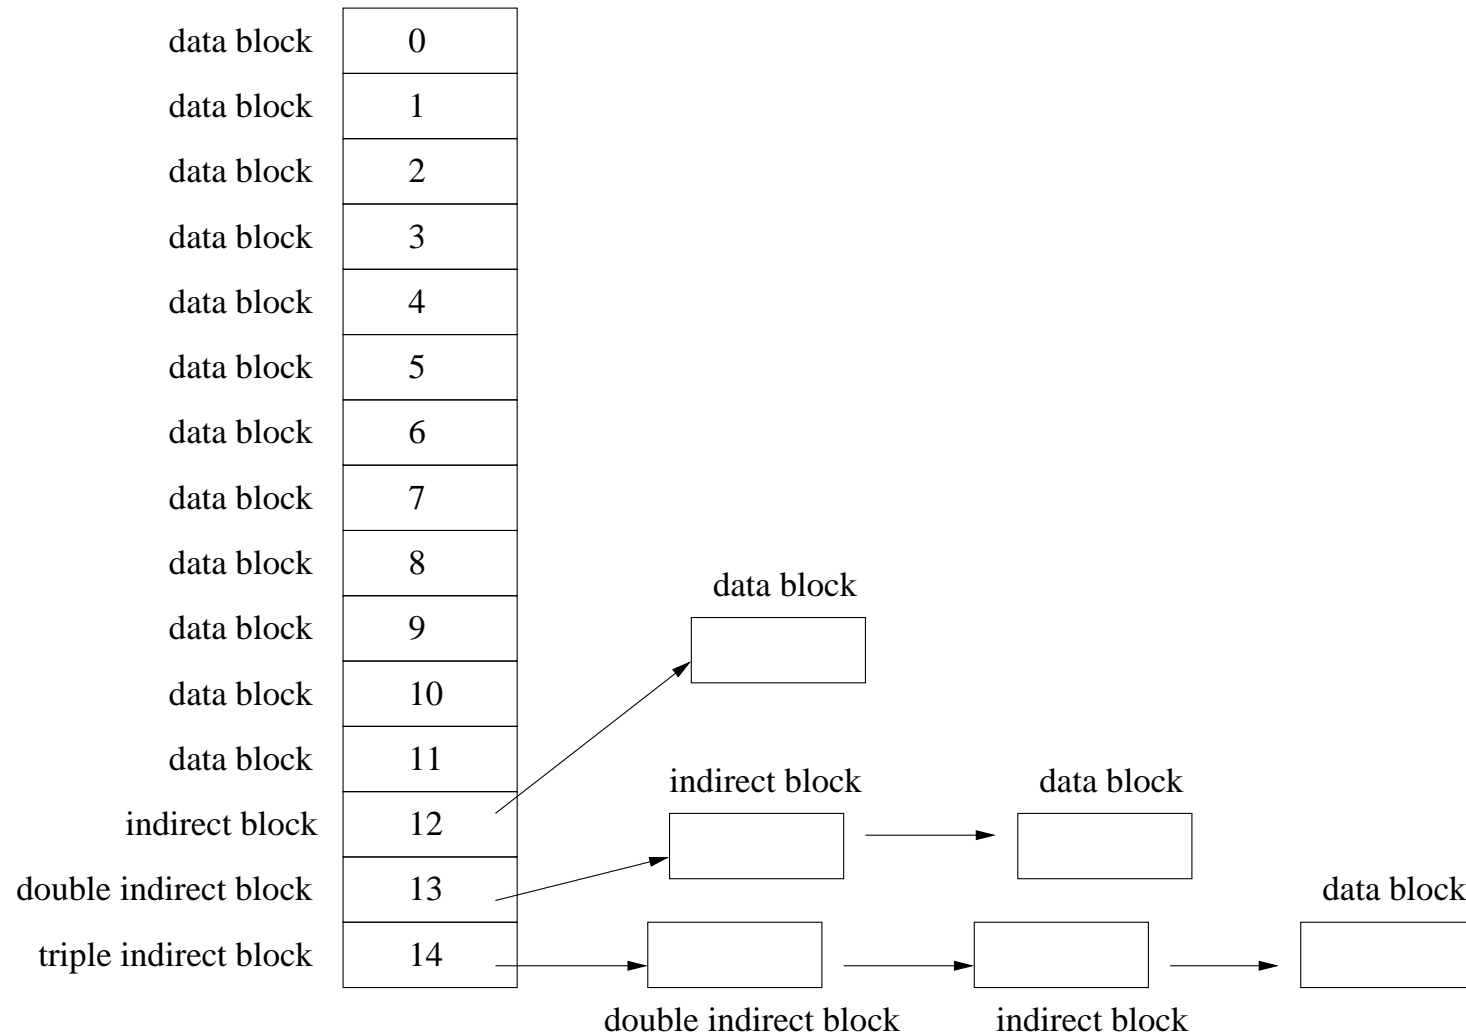


Structure of i-node

i-node in Unix file systems contains the following information

- owners
- timestamps
- size
- count:
- pointers to data blocks

block allocation of the Unix file system



block allocation of the Unix file sytem

If the block number is 4 bytes long, and each block is 4k bytes.

- datablock

$$12 \times 4k = 48k$$

- Single Indirect Blocks

$$1k \times 4kb = 4MB \text{ total space}$$

- Double Indirect Blocks


$$1 \times 1k \times 1k \times 4kb = 4GB \text{ total space}$$

- Triple Indirect Blocks

$$1 \times 1k \times 1k \times 1k \times 4kb = 4TB \text{ of total storage space}$$

block allocation of the Unix file system

How to record data block allocation for a file with fourteen disk blocks of content:




To list the inode numbers for all files under a directory, use


```
ls -laR directoryName
```

block allocation of the Unix file sytem

user view



system view



System calls for directories

mkdir, rmdir, unlink, link, rename, chdir

`$man 2 mkdir`

MKDIR(2)

Linux Programmers Manual

MKDIR(2)

NAME

`mkdir` - create a directory

SYNOPSIS

```
#include <sys/stat.h>
#include <sys/types.h>
```

```
int mkdir(const char *pathname, mode_t mode);
```

RETURN VALUE

`mkdir` returns zero on success, or `-1` if an error occurred (in which case, `errno` is set appropriately).

System calls for directories

```
$man 2 rmdir
```

RMDIR(2)

Linux Programmers Manual

RMDIR(2)

NAME

rmdir - delete a directory

SYNOPSIS

```
#include <unistd.h>
```

```
int rmdir(const char *pathname);
```

RETURN VALUE

On success, zero is returned. On error, -1 is returned, and `errno` is set appropriately.

System calls for directories

`$man 2 unlink`

UNLINK(2)

Linux Programmers Manual

UNLINK(2)

NAME

`unlink` - delete a name and possibly the file it refers to

SYNOPSIS

`#include <unistd.h>`

`int unlink(const char *pathname);`

RETURN VALUE

On success, zero is returned. On error, -1 is returned, and `errno` is set appropriately.

System calls for directories

`$man 2 link`

LINK(2)

Linux Programmers Manual

LINK(2)

NAME

`link` - make a new name for a file

SYNOPSIS

`#include <unistd.h>`

`int link(const char *oldpath, const char *newpath);`

RETURN VALUE

On success, zero is returned. On error, -1 is returned, and `errno` is set appropriately.

System calls for directories

```
$man 2 rename
```

```
RENAME(2)
```

```
Linux Programmers Manual
```

```
RENAME(2)
```

```
NAME
```

```
rename - change the name or location of a file
```

```
SYNOPSIS
```

```
#include <stdio.h>
```

```
int rename(const char *oldpath, const char *newpath);
```

System calls for directories

`$man 2 chdir`

CHDIR(2)

Linux Programmers Manual

CHDIR(2)

NAME

`chdir`, `fchdir` - change working directory

SYNOPSIS

`#include <unistd.h>`

`int chdir(const char *path);`

`int fchdir(int fd);`

RETURN VALUE

On success, zero is returned. On error, -1 is returned, and `errno` is set appropriately.

How pwd works

- Step 1: Note the inode number for “.”, call it n (use stat).
- Step 2: `chdir ..` (use `chdir`).
- Step 3: Find the name of the link with inode n (use `opendir`, `readdir`, `closedir`).
- Repeat step 2 and step 3 until you reach the top of the tree.

How pwd works

```
/* spwd.c: a simplified version of pwd
 *
 * starts in current directory and recursively
 * climbs up to root of filesystem, prints top part
 * then prints current part
 *
 * uses readdir() to get info about each thing
 *
 *      bug: prints an empty string if run from "/"
 **/
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <dirent.h>

ino_t get_inode(char *);
void    printpathto(ino_t);
void    inum_to_name(ino_t , char *, int );

int main()
{
    printpathto( get_inode( "." ) );           /* print path to here */
    putchar('\n');                             /* then add newline */
}
```

```
        return 0;
    }

void printpathto( ino_t this_inode )
/*
 *   prints path leading down to an object with this inode
 *   kindof recursive
 */
{
    ino_t my_inode ;
    char its_name[BUFSIZ];

    if ( get_inode("..") != this_inode )
    {
        chdir( ".." );                                /* up one dir */

        inum_to_name(this_inode,its_name,BUFSIZ);/* get its name*/

        my_inode = get_inode( "." );                  /* print head */
        printpathto( my_inode );                       /* recursively */
        printf("/%s", its_name );                     /* now print */
                                                    /* name of this */
    }
}
```



```
void inum_to_name(ino_t inode_to_find , char *namebuf, int buflen)
/*
 *   looks through current directory for a file with this inode
 *   number and copies its name into namebuf
 */
{
    DIR          *dir_ptr;                /* the directory */
    struct dirent *direntp;                /* each entry */

    dir_ptr = opendir( "." );
    if ( dir_ptr == NULL ){
        perror( "." );
        exit(1);
    }

    /*
     * search directory for a file with specified inum
     */

    while ( ( direntp = readdir( dir_ptr ) ) != NULL )
        if ( direntp->d_ino == inode_to_find )
        {
            strncpy( namebuf, direntp->d_name, buflen);
            namebuf[buflen-1] = '\0';    /* just in case */
        }
    }
```

```
        closedir( dir_ptr );
        return;
    }
    fprintf(stderr, "error looking for inum %d\n", inode_to_find);
    exit(1);
}


ino_t get_inode( char *fname )
/*
 *   returns inode number of the file
 */
{
    struct stat info;

    if ( stat( fname , &info ) == -1 ){
        fprintf(stderr, "Cannot stat ");
        perror(fname);
        exit(1);
    }
    return info.st_ino;
}
```

Multiple file systems: A tree of tree


user view: one tree

system view: two disks



Mount a one file system to a mount point of another file system

\$ mount options devicename mountpoint



Hard link and symbolic link

- Due to the fact that files in different file systems may have the same inode number, therefore hard link can only point to inodes in other file systems.
- symbolic link refers a file by name, not by inode number, therefore can point to a file in a different file system.

```
[wch@localhost chinese_story]$ ln /home/wch/software/jgrasp183.zip jgrap183.zip
ln: creating hard link 'jgrap183.zip' to '/home/wch/software/jgrasp183.zip':
Invalid cross-device link
```

```
[wch@localhost software]$ ls -l
total 5260
drwxr-xr-x  9 root root    4096 Nov 10 17:19 jdk1.5.0_06
drwxr-xr-x 10 root root    4096 Aug 17  2005 jgrasp
-rw-rw-r--  2 wch  wch   2669095 Jan  5 20:52 jgrasp183_link.zip
lrwxrwxrwx  1 wch  users      13 Mar 13 11:45 jgrasp183_slink.zip -> jgrasp183.zip
-rw-rw-r--  2 wch  wch   2669095 Jan  5 20:52 jgrasp183.zip
drwxrwxr-x 16 wch  wch    4096 Oct 21 21:28 matlab
```